

01_visualizaciones_exploratorias

November 20, 2025

1 Visualizaciones Exploratorias

Este notebook se enfoca en las seis visualizaciones requeridas para la Parte I del TP.

```
[14]: from pathlib import Path
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

DATA_DIR = Path('../data/nlp-getting-started')
TRAIN_PATH = DATA_DIR / 'train.csv'
sns.set_theme(style='whitegrid')
```

```
[15]: train_df = pd.read_csv(TRAIN_PATH)
train_df.head(30)
```

```
[15]:   id keyword location          text \
0    1      NaN     NaN  Our Deeds are the Reason of this #earthquake M...
1    4      NaN     NaN           Forest fire near La Ronge Sask. Canada
2    5      NaN     NaN  All residents asked to 'shelter in place' are ...
3    6      NaN     NaN  13,000 people receive #wildfires evacuation or...
4    7      NaN     NaN  Just got sent this photo from Ruby #Alaska as ...
5    8      NaN     NaN  #RockyFire Update => California Hwy. 20 closed...
6   10      NaN     NaN  #flood #disaster Heavy rain causes flash flood...
7   13      NaN     NaN  I'm on top of the hill and I can see a fire in...
8   14      NaN     NaN  There's an emergency evacuation happening now ...
9   15      NaN     NaN  I'm afraid that the tornado is coming to our a...
10  16      NaN     NaN  Three people died from the heat wave so far
11  17      NaN     NaN  Haha South Tampa is getting flooded hah- WAIT ...
12  18      NaN     NaN  #raining #flooding #Florida #TampaBay #Tampa 1...
13  19      NaN     NaN  #Flood in Bago Myanmar #We arrived Bago
14  20      NaN     NaN  Damage to school bus on 80 in multi car crash ...
15  23      NaN     NaN  What's up man?
16  24      NaN     NaN  I love fruits
17  25      NaN     NaN  Summer is lovely
18  26      NaN     NaN  My car is so fast
19  28      NaN     NaN  What a goooooooooooooal!!!!!!
```


20	31	NaN	NaN	this is ridiculous...
21	32	NaN	NaN	London is cool ;)
22	33	NaN	NaN	Love skiing
23	34	NaN	NaN	What a wonderful day!
24	36	NaN	NaN	L000000L
25	37	NaN	NaN	No way...I can't eat that shit
26	38	NaN	NaN	Was in NYC last week!
27	39	NaN	NaN	Love my girlfriend
28	40	NaN	NaN	Coooool :)
29	41	NaN	NaN	Do you like pasta?

	target
0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1
15	0
16	0
17	0
18	0
19	0
20	0
21	0
22	0
23	0
24	0
25	0
26	0
27	0
28	0
29	0

1.1 Visualización 1: Distribución del target

```
[16]: labels_map = {0: 'Falso (0)', 1: 'Desastre real (1)'}
class_counts = train_df['target'].value_counts().sort_index()
class_percentages = class_counts / class_counts.sum() * 100
plot_df = pd.DataFrame(
    {
        'label': [labels_map[idx] for idx in class_counts.index],
        'count': class_counts.values,
        'percentage': class_percentages.values
    }
)

fig, ax = plt.subplots(figsize=(6, 5))
bars = ax.bar(plot_df['label'], plot_df['count'], color=['#f58518', '#1098ad'])
for bar, pct in zip(bars, plot_df['percentage']):
    ax.text(
        bar.get_x() + bar.get_width() / 2,
        bar.get_height() + 50,
        f"{pct:.1f}%",
        ha='center',
        va='bottom',
        fontsize=11,
        color='#333333'
    )
ax.set_title('Distribución del target en el set de entrenamiento', fontsize=14)
ax.set_ylabel('Cantidad de tweets')
ax.set_xlabel('Clase')
plt.tight_layout()
plt.show()
```



Okey más o menos igualmente distribuidos los reales de los Falsos. Algunos más falsos que reales.

1.2 Visualización 2: Impacto del conteo de palabras en la clasificación


```
[17]: word_counts = train_df['text'].fillna('').str.split().str.len()
plot_df = pd.DataFrame({
    'Clase': train_df['target'].map({0: 'Falso (0)', 1: 'Desastre real (1)'}),
    'Cantidad de palabras': word_counts
})

fig, ax = plt.subplots(figsize=(7, 5))
sns.boxplot(
    data=plot_df,
    x='Clase',
    y='Cantidad de palabras',
    hue='Clase',
    palette=['#f58518', '#1098ad'],
    ax=ax
```

```

)
legend = ax.legend_
if legend is not None:
    legend.remove()
ax.set_title('Cantidad de palabras por clase', fontsize=14)
ax.set_xlabel('Clase')
ax.set_ylabel('Cantidad de palabras')
plt.tight_layout()
plt.show()

```



Ok, parece que no hay mucha diferencia en la cantidad de palabras entre tweets de desastre y no desastre. Podríamos intentar hacer un análisis más detallado, como cantidad de letras por tweet.

1.3 Visualización 3: Ubicación geográfica de los tweets

Voy a graficar la proporción de tweets de desastre por ubicación con alta cobertura para evaluar el contexto geográfico.

```
[18]: location_clean = train_df['location'].fillna('Unknown').str.strip()
location_counts = location_clean.value_counts()
top_locations = location_counts[location_counts >= 40].head(10).index
plot_df = (


```

```

train_df.assign(location=location_clean)
.loc[lambda df: df['location'].isin(top_locations)]
.groupby('location')['target']
.agg(tweets='size', disaster_rate='mean')
.sort_values('disaster_rate', ascending=False)
.reset_index()
)
plot_df['disaster_rate_pct'] = plot_df['disaster_rate'] * 100

fig, ax = plt.subplots(figsize=(9, 5))
sns.barplot(
    data=plot_df,
    x='disaster_rate_pct',
    y='location',
    hue='location',
    palette=sns.color_palette('viridis', len(plot_df)),
    dodge=False,
    ax=ax
)
legend = ax.legend_
if legend is not None:
    legend.remove()
for idx, row in plot_df.iterrows():
    ax.text(
        row['disaster_rate_pct'] + 1,
        idx,
        f'{row["disaster_rate_pct"]:.1f} ({row["tweets"]}) tweets',
        va='center',
        fontsize=10,
        color="#333333"
    )
ax.set_title('Porcentaje de tweets de desastre real por ubicaciones',□
             fontsize=14)
ax.set_xlabel('Tweets clasificados como desastre (%)')
ax.set_ylabel('Ubicación')
plt.tight_layout()
plt.show()

```



```
[19]: location_clean = train_df['location'].fillna('Unknown').str.strip()
all_locations = location_clean.value_counts()

print(f"Total unique locations: {len(all_locations)}")
print(f"\nTop 50 locations by count:")
print("-"*60)
for loc, count in all_locations.head(50).items():
    print(f"{loc:40s} : {count:4d} tweets")

print(f"\n\nLocations with 'USA' or 'United States' or 'America':")
print("-"*60)
usa_related = all_locations[all_locations.index.str.contains('USA|United
    -States|America|U.S', case=False, na=False)]
for loc, count in usa_related.items():
    print(f"{loc:40s} : {count:4d} tweets")
```

Total unique locations: 3279

Top 50 locations by count:

Unknown	:	2536 tweets
USA	:	105 tweets
New York	:	73 tweets
United States	:	50 tweets
London	:	46 tweets
Nigeria	:	32 tweets
Canada	:	30 tweets

UK	:	27 tweets
Los Angeles, CA	:	26 tweets
India	:	24 tweets
Mumbai	:	22 tweets
Washington, DC	:	21 tweets
California	:	20 tweets
Kenya	:	20 tweets
Chicago, IL	:	19 tweets
Worldwide	:	19 tweets
Australia	:	19 tweets
New York, NY	:	16 tweets
Los Angeles	:	15 tweets
California, USA	:	15 tweets
Everywhere	:	15 tweets
Indonesia	:	14 tweets
Florida	:	14 tweets
Washington, D.C.	:	14 tweets
United Kingdom	:	14 tweets
San Francisco	:	14 tweets
Toronto	:	12 tweets
Ireland	:	12 tweets
NYC	:	12 tweets
Earth	:	12 tweets
Seattle	:	11 tweets
Chicago	:	11 tweets
San Francisco, CA	:	11 tweets
Texas	:	11 tweets
London, UK	:	10 tweets
London, England	:	10 tweets
Atlanta, GA	:	10 tweets
Dallas, TX	:	10 tweets
ss	:	10 tweets
New York City	:	10 tweets
Sacramento, CA	:	10 tweets
Nashville, TN	:	9 tweets
US	:	9 tweets
Manchester	:	9 tweets
World	:	9 tweets
Denver, Colorado	:	9 tweets
Houston, TX	:	9 tweets
San Diego, CA	:	9 tweets
Scotland	:	9 tweets
304	:	9 tweets

Locations with 'USA' or 'United States' or 'America':

=====		
USA	:	105 tweets

United States	:	50 tweets
California, USA	:	15 tweets
Pennsylvania, USA	:	7 tweets
California, United States	:	6 tweets
Texas, USA	:	5 tweets
Florida, USA	:	5 tweets
New York, USA	:	5 tweets
U.S.A	:	4 tweets
North Carolina, USA	:	4 tweets
Massachusetts, USA	:	4 tweets
Tulsa, Oklahoma	:	3 tweets
Oregon, USA	:	3 tweets
Upstairs.	:	3 tweets
New Orleans ,Louisiana	:	3 tweets
St. Louis, MO	:	3 tweets
America of Founding Fathers	:	3 tweets
Hawaii, USA	:	3 tweets
Illinois, USA	:	3 tweets
America	:	3 tweets
Virginia, USA	:	3 tweets
South, USA	:	3 tweets
Ohio, USA	:	2 tweets
Washington, USA	:	2 tweets
somewhere USA	:	2 tweets
Moncton, New Brunswick	:	2 tweets
Georgia, USA	:	2 tweets
USA (Formerly @usNOAAgov)	:	2 tweets
Louisville, KY	:	2 tweets
Michigan, USA	:	2 tweets
West Virginia, USA	:	2 tweets
Redding, California, USA	:	2 tweets
Jerusalem	:	2 tweets
Louisiana	:	2 tweets
United States of America	:	2 tweets
vancouver usa	:	2 tweets
The American Wasteland (MV)	:	2 tweets
Tucson, AZ	:	1 tweets
Americas Newsroom	:	1 tweets
New Mexico, USA	:	1 tweets
Outside The Matrix, I Think.	:	1 tweets
Missouri, USA	:	1 tweets
somewhere outside	:	1 tweets
Tucson, Arizona	:	1 tweets
Las Vegas, NV USA	:	1 tweets
3?3?7?SLOPelousas??2?2?5?	:	1 tweets
Afghanistan, USA	:	1 tweets
Littleton, CO, USA	:	1 tweets
Virginia, United States	:	1 tweets

Northern Kentucky, USA	:	1 tweets
America New Zealand	:	1 tweets
New Brunswick, NJ	:	1 tweets
Moscow, Russia	:	1 tweets
In the clouds...	:	1 tweets
SEATTLE, WA USA	:	1 tweets
USA , AZ	:	1 tweets
Lancaster, Pennsylvania, USA	:	1 tweets
Lindenhurst	:	1 tweets
Cape Cod, Massachusetts USA	:	1 tweets
Purgatory, USA	:	1 tweets
Sunshine Coast, Queensland	:	1 tweets
USA, WA	:	1 tweets
USAoV	:	1 tweets
U.S	:	1 tweets
Made in America	:	1 tweets
Los Angeles,CA, USA	:	1 tweets
the void, U.S.A	:	1 tweets
Detroit, MI, United States	:	1 tweets
Maryland, USA	:	1 tweets
Texas-USA ¢ ?	:	1 tweets
Washington, Krasnodar (Russia)	:	1 tweets
Basketball City, USA	:	1 tweets
Iowa, USA	:	1 tweets
USA/SO FLORIDA via BROOKLYN NY	:	1 tweets
St. Louis Mo.	:	1 tweets
Hickville, USA	:	1 tweets
U.S.	:	1 tweets
USA, Haiti, Nepal	:	1 tweets
Philadelphia, PA USA	:	1 tweets
U.S.A. - Global Members Site	:	1 tweets
Ù¢5 Ù¢12 Ù¢14 Ù¢ åÈ#SaviourSquadåÇ	:	1 tweets
Saint Louis, Missouri	:	1 tweets
Jonesboro, Arkansas USA	:	1 tweets
Hudson Valley, NY	:	1 tweets
Alaska, USA	:	1 tweets
Nevada, USA	:	1 tweets
Alexandria, VA, USA	:	1 tweets
U.S. Northern Virginia	:	1 tweets
Jerusalem!	:	1 tweets
USA, Alabama	:	1 tweets
New Orleans, Louisiana	:	1 tweets
Northern California U.S.A.	:	1 tweets
Eureka, California, USA	:	1 tweets
Los Angeles... CA... USA	:	1 tweets
AKRON OHIO USA	:	1 tweets
Livingston, IL U.S.A.	:	1 tweets
The Hammock, FL, USA	:	1 tweets

Jerusalem, Israel	:	1 tweets
United States where it's warm	:	1 tweets
Midwestern USA	:	1 tweets
louisville, kentucky	:	1 tweets
St. Louis, Missouri	:	1 tweets
Tucson, Az	:	1 tweets
Freeport IL. USA	:	1 tweets
Hawaii USA	:	1 tweets
I ACCEPT SONG REQUESTS	:	1 tweets
North America	:	1 tweets
Elk Grove, CA, USA	:	1 tweets
Books Published, USA	:	1 tweets
Orlando,FL USA	:	1 tweets
Georgia, U.S.A.	:	1 tweets
Madison, WI & St. Louis MO	:	1 tweets
Russia	:	1 tweets
San Jose, CA, USA	:	1 tweets
st.louis county missouri	:	1 tweets
Oklahoma, USA	:	1 tweets
St. Louis, Mo	:	1 tweets
Tennessee, USA	:	1 tweets
Wausau, Wisconsin	:	1 tweets
Jonesboro, AR MO, IOWA USA	:	1 tweets
sí£o luis	:	1 tweets
Colorado, USA	:	1 tweets
Phoenix, Arizona, USA	:	1 tweets
Malibu/SantaFe/Winning!	:	1 tweets
USA, North Dakota	:	1 tweets
San Luis Obispo, CA	:	1 tweets
828/704(Soufside)/while looking goofy in NJ	:	1 tweets
Pro-American and Anti-#Occupy	:	1 tweets
Not a U.S resident	:	1 tweets
MI,USA	:	1 tweets
missouri USA	:	1 tweets
New York, United States	:	1 tweets
russia	:	1 tweets
LP, MN USA	:	1 tweets
USA - Canada - Europe - Asia	:	1 tweets
Tulsa, OK	:	1 tweets
Upstate New York	:	1 tweets
U.S.A. FEMA Region 5	:	1 tweets
Spain - China - Latin America.	:	1 tweets
Philadelphia, Pennsylvania USA	:	1 tweets
trapped in America	:	1 tweets
Montana, USA	:	1 tweets
Utah, USA	:	1 tweets
PA.USA	:	1 tweets
Madison, Wisconsin, USA	:	1 tweets

Nomad, USA	:	1 tweets
North East Unsigned Radio	:	1 tweets
Duval, WV 25573, USA ?	:	1 tweets
Soufside	:	1 tweets
PSA Nursing	:	1 tweets
Louisiana, USA	:	1 tweets
Bon Temps Louisiana	:	1 tweets
West Coast, USA	:	1 tweets
U.S.A and Canada	:	1 tweets
Minority Privilege, USA	:	1 tweets
New Jersey, USA	:	1 tweets
Reston, VA, USA	:	1 tweets
Tornado Alley, USA	:	1 tweets
PA, USA	:	1 tweets
LITTLETON, CO, USA, TERRAN	:	1 tweets
South Carolina, USA	:	1 tweets
Atlanta, Georgia USA	:	1 tweets
Kentucky, USA	:	1 tweets
Gotham City,USA	:	1 tweets
Cambridge, Massachusetts, U.S.	:	1 tweets
St Louis, MO	:	1 tweets
Wisconsin, USA	:	1 tweets
eBooks, North America	:	1 tweets
Indiana, USA	:	1 tweets
Very SW CA, USA...Draenor	:	1 tweets
New Hampshire, USA	:	1 tweets
West Coast, Cali USA	:	1 tweets
Savage States of America	:	1 tweets
GrC Founder, 8,000 Subscribers	:	1 tweets
CT, USA	:	1 tweets
Proud @BuckMasonUSA supporter!	:	1 tweets
N. California USA	:	1 tweets
Vermont, USA	:	1 tweets
pettyville, usa	:	1 tweets
Florida USA	:	1 tweets
St. Louis	:	1 tweets
North East USA	:	1 tweets
Alabama, USA	:	1 tweets

claro, el campo location está totalmente destruido. Necesito armarme una feature “country”. Usé la API Pública de geocoders para sacar el país de la ubicación dada. Dejo el código abajo!

[20]: `import json`

`'''`

Usé este script para generar `location_to_country.json` que me mapea las ubicaciones de los usuarios a países.

```

from geopy.geocoders import Nominatim
from time import sleep

geolocator = Nominatim(user_agent="tweet-checker")

unique_locations = train_df['location'].dropna().unique()
location_to_country = {}

print(f"Processing {len(unique_locations)} unique locations...")

for i, location in enumerate(unique_locations):
    if pd.isna(location) or location.strip() == '' or location == 'Unknown':
        location_to_country[location] = None
        continue

    try:
        result = geolocator.geocode(location, addressdetails=True, timeout=10)
        if result and 'address' in result.raw:
            country = result.raw['address'].get('country')
            location_to_country[location] = country
        else:
            location_to_country[location] = None
    except Exception as e:
        location_to_country[location] = None

    if (i + 1) % 50 == 0:
        print(f"Processed {i + 1}/{len(unique_locations)}")

    sleep(1)

print(f"Done. Detected {sum(1 for v in location_to_country.values() if v is not None)} countries")

for loc, country in list(location_to_country.items())[:20]:
    print(f"{loc}: {country}")

with open('../data/location_to_country.json', 'w', encoding='utf-8') as f:
    json.dump(location_to_country, f, ensure_ascii=False, indent=2)

'''

with open('../data/location_to_country.json', 'r', encoding='utf-8') as f:
    location_to_country = json.load(f)

```

```
[21]: train_df['country'] = train_df['location'].map(location_to_country)

print(f"Rows with country: {train_df['country'].notna().sum()}")
print(f"Rows without country: {train_df['country'].isna().sum()}")

print("\nTop countries:")
print(train_df['country'].value_counts().head(15))

train_df[train_df['country'].notna()][['location', 'country', 'text']].head(10)
```

Rows with country: 4060

Rows without country: 3553

Top countries:

country	
United States of America	2153
United Kingdom	438
Canada	222
Australia	150
India	138
Nigeria	76
Philippines	45
Kenya	40
South Africa	37
France	37
Indonesia	33
Brasil	30
	26
Deutschland	26
Éire / Ireland	26

Name: count, dtype: int64

```
[21]:          location           country \
31      Birmingham           United Kingdom
34  Philadelphia, PA  United States of America
35      London, UK           United Kingdom
36      Pretoria            South Africa
37      World Wide!!  United States of America
39      Paranaque City       Philippines
42      milky way             Sesel
46  GREENSBORO,NORTH CAROLINA  United States of America
49      England.           United Kingdom
50  Sheffield Township, Ohio  United States of America

                                         text
31  @bbcmtd Wholesale Markets ablaze http://t.co/l...
34      Crying out for more! Set me ablaze
35  On plus side LOOK AT THE SKY LAST NIGHT IT WAS...
```

```

36 @PhDSquares #mufc they've built so much hype a...
37 INEC Office in Abia Set Ablaze - http://t.co/3...
39                                     Ablaze for you Lord :D
42 Had an awesome time visiting the CFC head offi...
46 How the West was burned: Thousands of wildfire...
49 First night with retainers in. It's quite weir...
50 Deputies: Man shot before Brighton home set ab...

```


Me cierra. El que ni idea es milky way pero no había mucha esperanza para esa.

```
[22]: country_counts = train_df['country'].value_counts()
top_countries = country_counts[country_counts >= 40].head(10).index

plot_df = (
    train_df[train_df['country'].isin(top_countries)]
    .groupby('country')['target']
    .agg(tweets='size', disaster_rate='mean')
    .sort_values('disaster_rate', ascending=False)
    .reset_index()
)
plot_df['disaster_rate_pct'] = plot_df['disaster_rate'] * 100

fig, ax = plt.subplots(figsize=(9, 5))
sns.barplot(
    data=plot_df,
    x='disaster_rate_pct',
    y='country',
    hue='country',
    palette=sns.color_palette('viridis', len(plot_df)),
    dodge=False,
    ax=ax
)
legend = ax.legend_
if legend is not None:
    legend.remove()
for idx, row in plot_df.iterrows():
    ax.text(
        row['disaster_rate_pct'] + 1,
        idx,
        f'{row["disaster_rate_pct"]:.1f}% ({row["tweets"]}) tweets",
        va='center',
        fontsize=10,
        color="#333333"
    )
ax.set_title('Porcentaje de tweets de desastre real por país', fontsize=14)
ax.set_xlabel('Tweets clasificados como desastre (%)')
ax.set_ylabel('País')
```

```
plt.tight_layout()
plt.show()
```



Okey!! Vemos data bastante más útil ahora, la normalización por país diferenció bastante más! Aunque claro, se mencionó bastante en una clase de la materia, que ante menos población, la variabilidad es mucho mayor. Aun así, me encantaron los resultados, me quedo fuerte con esta feature.

1.4 Visualización 4: Longitud de texto vs target

```
[23]: text_lengths = train_df['text'].fillna('').str.len()
plot_df = pd.DataFrame({
    'Clase': train_df['target'].map({0: 'Falso (0)', 1: 'Desastre real (1)'}) ,
    'Longitud de texto (caracteres)': text_lengths
})

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

sns.boxplot(
    data=plot_df,
    x='Clase',
    y='Longitud de texto (caracteres)',
    hue='Clase',
    palette=['#f58518', '#1098ad'],
    ax=ax1
)
legend = ax1.legend_
```


```

if legend is not None:
    legend.remove()
ax1.set_title('Distribución de longitud de texto por clase', fontsize=14)
ax1.set_xlabel('Clase')
ax1.set_ylabel('Longitud (caracteres)')

sns.kdeplot(
    data=plot_df,
    x='Longitud de texto (caracteres)',
    hue='Clase',
    fill=True,
    palette=['#f58518', '#1098ad'],
    alpha=0.5,
    ax=ax2
)
ax2.set_title('Densidad de longitud de texto por clase', fontsize=14)
ax2.set_xlabel('Longitud (caracteres)')
ax2.set_ylabel('Densidad')

plt.tight_layout()
plt.show()

```



okey! Sorprendentemente los tweets falsos tienen menos texto que los otros. inesperado pero bueno.

1.5 Visualización 5: Entidades y símbolos

Resaltar cómo la presencia de URLs, hashtags o mentions se relaciona con el target.

```
[24]: train_df['has_url'] = train_df['text'].fillna('').str.contains(r'http[s]?://', regex=True)
train_df['has_hashtag'] = train_df['text'].fillna('').str.contains(r'#\w+', regex=True)
```

```

train_df['has_mention'] = train_df['text'].fillna('').str.contains(r'@\w+', ↴
    ↴regex=True)

features = ['has_url', 'has_hashtag', 'has_mention']
feature_names = ['URLs', 'Hashtags', 'Mentions']

stats = []
for target_val in [0, 1]:
    for feature, name in zip(features, feature_names):
        pct = train_df[train_df['target'] == target_val][feature].mean() * 100
        stats.append({
            'Clase': 'Falso (0)' if target_val == 0 else 'Desastre real (1)',
            'Feature': name,
            'Porcentaje': pct
        })
plot_df = pd.DataFrame(stats)

fig, ax = plt.subplots(figsize=(10, 6))
x = np.arange(len(feature_names))
width = 0.35

false_data = plot_df[plot_df['Clase'] == 'Falso (0)']['Porcentaje'].values
disaster_data = plot_df[plot_df['Clase'] == 'Desastre real (1)']['Porcentaje'].values
    ↴values

bars1 = ax.bar(x - width/2, false_data, width, label='Falso (0)', ↴
    ↴color="#f58518")
bars2 = ax.bar(x + width/2, disaster_data, width, label='Desastre real (1)', ↴
    ↴color="#1098ad")


for bars in [bars1, bars2]:
    for bar in bars:
        height = bar.get_height()
        ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{height:.1f}%', ↴
            ha='center', va='bottom', fontsize=10)

ax.set_xlabel('Tipo de entidad/símbolo')
ax.set_ylabel('Porcentaje de tweets')
ax.set_title('Presencia de URLs, hashtags y mentions por clase', fontsize=14)
ax.set_xticks(x)
ax.set_xticklabels(feature_names)
ax.legend()
ax.set_ylim(0, max(plot_df['Porcentaje']) * 1.15)

plt.tight_layout()

```

```
plt.show()
```



Bueno, por lo menos hay una correspondencia entre simbolos y tweets de desastre. Nuevamente pensé que sería al revés la relación, pero lo tomo. Parece que el más determinante es si tiene o no URL.

1.6 Visualización 6: Nube de palabras - Palabras asociadas con desastres

Nube de palabras mostrando las palabras más frecuentes (mín 5 tweets). Tamaño = frecuencia, Color = tasa de desastre.

```
[25]: from wordcloud import WordCloud
import re
from collections import Counter

# Filtro un par de palabras que no me interesa ver distribuciones: (igual me quedé cortísimo)
stop_words = set(['the', 'a', 'an', 'and', 'or', 'but', 'in', 'on', 'at', 'to', 'for',
                  'of', 'with', 'by', 'from', 'as', 'is', 'was', 'are', 'been', 'be',
                  'have', 'has', 'had', 'do', 'does', 'did', 'will', 'would', 'should',
                  'could', 'may', 'might', 'must', 'can', 'this', 'that', 'these', 'those',
```

```

'i', 'you', 'he', 'she', 'it', 'we', 'they', 'what', 'which', ↵
↳ 'who',
'when', 'where', 'why', 'how', 'all', 'each', 'every', ↵
↳ 'both', 'few',
'more', 'most', 'other', 'some', 'such', 'no', 'nor', 'not', ↵
↳ 'only',
'own', 'same', 'so', 'than', 'too', 'very', 'just', 'https', ↵
↳ 'http',
'co', 't', 's', 'm', 're', 've', 'll', 'd'])

word_stats = []
for idx, row in train_df.iterrows():
    text = str(row['text']).lower()
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'[^a-zA-Z\s]', ' ', text)
    words = text.split()
    words = [w for w in words if w not in stop_words and len(w) > 2]

    for word in set(words):
        if word not in word_stats:
            word_stats[word] = {'count': 0, 'disaster_count': 0}
            word_stats[word]['count'] += 1
            word_stats[word]['disaster_count'] += row['target']

word_data = []
for word, stats in word_stats.items():
    if stats['count'] >= 5:
        disaster_rate = stats['disaster_count'] / stats['count']
        word_data.append({
            'word': word,
            'count': stats['count'],
            'disaster_rate': disaster_rate
        })


word_df = pd.DataFrame(word_data).sort_values('count', ascending=False)

frequencies = {row['word']: row['count'] for _, row in word_df.iterrows()}
disaster_rates = {row['word']: row['disaster_rate'] for _, row in word_df.
↳ iterrows()}

def color_func(word, **kwargs):
    rate = disaster_rates.get(word, 0.5)
    if rate > 0.7:
        return 'rgb(200, 50, 50)'
    elif rate > 0.5:
        return 'rgb(255, 150, 50)'
    elif rate > 0.3:

```

```
        return 'rgb(100, 150, 200)'  
    else:  
        return 'rgb(50, 100, 200)'  
  
wc = WordCloud(  
    width=1200,  
    height=600,  
    background_color='white',  
    max_words=100,  
    relative_scaling=0.5,  
    min_font_size=10  
).generate_from_frequencies(frequencies)  
  
fig, ax = plt.subplots(figsize=(16, 8))  
ax.imshow(wc.recolor(color_func=color_func), interpolation='bilinear')  
ax.axis('off')  
ax.set_title('Word Cloud - Palabras asociadas con desastres\n(Tamaño =  
↳frecuencia, Color = disaster rate: Rojo >70%, Naranja >50%, Celeste >30%,  
↳Azul <30%)',  
            fontsize=14, pad=20)  
plt.tight_layout()  
plt.show()  
  
print(f"\nTop 20 palabras por disaster rate (min 5 tweets):")  
print(word_df.nlargest(20, 'disaster_rate')[['word', 'count', 'disaster_rate']])
```



Top 20 palabras por disaster rate (min 5 tweets):

	word	count	disaster_rate
843	northern	64	1.0
407	debris	49	1.0
1987	severe	44	1.0
1763	derailment	40	1.0
2044	legionnaires	39	1.0
2093	migrants	37	1.0
2348	investigators	37	1.0
926	mosque	33	1.0
2542	detonated	31	1.0
1577	pkk	31	1.0
1556	turkey	27	1.0
949	israeli	27	1.0
483	helicopter	25	1.0
2565	conclusively	25	1.0
2196	saipan	25	1.0
2356	projected	24	1.0
2051	signs	23	1.0
2384	trench	23	1.0
2484	refugio	22	1.0
2486	costlier	22	1.0

Okey hay palabras que claramente se asocian más a desastres, como “news”, “fire”, “accident”, etc. Mientras que otras como “like”, “get”, “your” se asocian más a tweets no relacionados con desastres. Igual son palabras bastante generales. filtré algunas con stop words pero en general no hay nada muy específico a los no desastres, pero bueno, si estás tweeteando un desastre probablemente no digas “like” o “your”. Debido a la seriedad del tema.

1.7 Visualización 7: Análisis de sentimiento del texto del tweet

```
[26]: from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer

analyzer = SentimentIntensityAnalyzer()

def get_sentiment_compound(text):
    if pd.isna(text) or text.strip() == '':
        return 0.0
    return analyzer.polarity_scores(text)['compound']

train_df['sentiment'] = train_df['text'].apply(get_sentiment_compound)

bins = [-1, -0.6, -0.2, 0.2, 0.6, 1]
labels = ['Muy negativo\n(-1 a -0.6)', 'Negativo\n(-0.6 a -0.2)', 'Neutral\n(-0.2 a 0.2)',
          'Positivo\n(0.2 a 0.6)', 'Muy positivo\n(0.6 a 1)']
train_df['sentiment_bin'] = pd.cut(train_df['sentiment'], bins=bins, labels=labels, include_lowest=True)
```

```

plot_df = train_df.groupby(['sentiment_bin', 'target']).size() .
    unstack(fill_value=0)

plot_df_pct = plot_df.div(plot_df.sum(axis=0), axis=1) * 100
plot_df_pct.columns = ['Falso (0)', 'Desastre real (1)']

fig, ax = plt.subplots(figsize=(12, 6))
plot_df_pct.plot(kind='bar', ax=ax, color=['#f58518', '#1098ad'], width=0.7)

for container in ax.containers:
    ax.bar_label(container, fmt='%.1f%%', padding=3)

ax.set_title('Distribución de sentiment por clase (normalizado)', fontsize=14)
ax.set_xlabel('Rango de sentiment (VADER compound score)')
ax.set_ylabel('Porcentaje de tweets (%)')
ax.legend(title='Clase', loc='upper right')
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
plt.tight_layout()
plt.show()

print("\nEstadísticas de sentiment por clase:")
print(train_df.groupby('target')['sentiment'].describe())


```

/tmp/ipykernel_76021/4106497693.py:17: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

plot_df = train_df.groupby(['sentiment_bin',
'target']).size().unstack(fill_value=0)

```



Estadísticas de sentiment por clase:

target	count	mean	std	min	25%	50%	75%	max
0	4342.0	-0.052444	0.465568	-0.9883	-0.4588	0.0000	0.3182	0.9730
1	3271.0	-0.267240	0.434629	-0.9686	-0.6390	-0.3182	0.0000	0.9471

Okey!!! Me gustó este, realmente parece haber correlación entre el sentimiento del tweet y si es desastre o no. Sobretodo se ve muy fuerte en las puntas. Los tweets de desastre tienden a tener sentimientos más negativos, mientras que los no desastres son más neutrales o positivos. Tiene mucho sentido así que voy a re usar la feature de sentimiento en el modelo.