

# 02\_regresion\_lineal

November 20, 2025

```
[1]: from pathlib import Path
import numpy as np
import pandas as pd
from sklearn.model_selection import StratifiedKFold, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score
from sklearn.feature_extraction.text import TfidfVectorizer
from scipy.stats import loguniform
```

## 1 Creación de Features Categóricas

### 1.0.1 Categórica 1: País del usuario

Agrego países calculados en la visualización 3 del notebook 01

```
[2]: import json

DATA_DIR = Path('../data/nlp-getting-started')
TRAIN_PATH = DATA_DIR / 'train.csv'
TEST_PATH = DATA_DIR / 'test.csv'
LOCATION_TO_COUNTRY_PATH = Path('../data/location_to_country.json')
RANDOM_SEED = 27

train_df = pd.read_csv(TRAIN_PATH)
test_df = pd.read_csv(TEST_PATH)

with open(LOCATION_TO_COUNTRY_PATH, 'r', encoding='utf-8') as f:
    location_to_country = json.load(f)

train_df['country'] = train_df['location'].map(location_to_country)
test_df['country'] = test_df['location'].map(location_to_country)

print(f"Train shape: {train_df.shape}")
print(f"Test shape: {test_df.shape}")
print(f"\nCountries in train: {train_df['country'].notna().sum()}/{len(train_df)}")
print(f"Countries in test: {test_df['country'].notna().sum()}/{len(test_df)}")
```

```
train_df.head()
```

```
Train shape: (7613, 6)
Test shape: (3263, 5)

Countries in train: 4060/7613
Countries in test: 882/3263
```

```
[2]:    id keyword location          text \
0     1      NaN      NaN  Our Deeds are the Reason of this #earthquake M...
1     4      NaN      NaN  Forest fire near La Ronge Sask. Canada
2     5      NaN      NaN  All residents asked to 'shelter in place' are ...
3     6      NaN      NaN  13,000 people receive #wildfires evacuation or...
4     7      NaN      NaN  Just got sent this photo from Ruby #Alaska as ...

      target country
0         1      NaN
1         1      NaN
2         1      NaN
3         1      NaN
4         1      NaN
```

Claro, es genial poder clasificar los que tienen location. Pero bueno, vemos acá que del test solo 882 tienen country. Nos sirve para alguna decisión de un arbol ponele, pero es limitado...

### 1.0.2 Categórica 2: Has\_Url

Vimos en la visu 5 que había fuerte diferencia en la proporción de tweets con url entre clases.

```
[3]: train_df['has_url'] = train_df['text'].fillna('').str.contains(r'http[s]?://',  
                     regex=True).astype(str)  
test_df['has_url'] = test_df['text'].fillna('').str.contains(r'http[s]?://',  
                     regex=True).astype(str)  
  
print("has_url distribution in train:")  
print(train_df['has_url'].value_counts())  
print(f"\nPercentage with URL: {train_df['has_url'].eq('True').mean()*100:.  
      1f}%")  
  
print("\nhas_url distribution in test:")  
print(test_df['has_url'].value_counts())  
print(f"\nPercentage with URL: {test_df['has_url'].eq('True').mean()*100:.1f}%")  
  
categorical_features = ['country', 'has_url']
```

```
has_url distribution in train:  
has_url  
True      3971  
False     3642
```

```
Name: count, dtype: int64
Percentage with URL: 52.2%
has_url distribution in test:
has_url
True      1731
False     1532
Name: count, dtype: int64
Percentage with URL: 53.0%
```

## 2 Features numéricas

### 2.0.1 Numérica 1: text\_length

La visu 4 mostró que la longitud del tweet es un indicador levemente útil, puede aportar algo.

```
[4]: train_df['text_length'] = train_df['text'].fillna('').str.len()
test_df['text_length'] = test_df['text'].fillna('').str.len()

print("Text length stats:")
print(f"Train - mean: {train_df['text_length'].mean():.1f}, std:{train_df['text_length'].std():.1f}")
print(f"Test - mean: {test_df['text_length'].mean():.1f}, std:{test_df['text_length'].std():.1f}")

numeric_features = ['text_length']
categorical_features = ['country', 'has_url']
embedding_feature = 'text'
```

```
Text length stats:
Train - mean: 101.0, std: 33.8
Test - mean: 102.1, std: 34.0
```

### 2.0.2 Numérica 2: sentiment\_score

Usé VADER para calcular un score de sentimiento del texto del tweet. La idea es que tweets con sentimientos muy marcados (positivos o negativos) pueden ser más probables de ser reales, mientras que los neutrales pueden ser más sospechosos.

```
[5]: from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer

analyzer = SentimentIntensityAnalyzer()

def get_sentiment(text):
    if pd.isna(text) or text.strip() == '':
        return 0.5
```

```
compound = analyzer.polarity_scores(text)['compound']
return (compound + 1) / 2

train_df['sentiment_score'] = train_df['text'].apply(get_sentiment)
test_df['sentiment_score'] = test_df['text'].apply(get_sentiment)

print("Sentiment score stats:")
print(f"Train - mean: {train_df['sentiment_score'].mean():.3f}, std:{train_df['sentiment_score'].std():.3f}")
print(f"Test - mean: {test_df['sentiment_score'].mean():.3f}, std:{test_df['sentiment_score'].std():.3f}")

print("\nSentiment score distribution in train:")
print(train_df['sentiment_score'].describe())

numeric_features = ['text_length', 'sentiment_score']
```

### Sentiment score stats:

Train - mean: 0.428, std: 0.232

Test - mean: 0.426, std: 0.229

Sentiment score distribution in train:

count 7613.00000

mean 0.427633

std 0.232416

min 0.005850

25% 0.228850

50% 0.500000

75% 0.538600

max 0.986500

Name: sentiment scor

<sup>1</sup> See also the discussion of the relationship between the two in the Introduction.

Voy a usar esta feature en el modelo y ver cómo sale.

```
[6]: train_df.head(10)
```

[6]:		id	keyword	location	text
0	1	NaN	NaN	Our Deeds are the Reason of this #earthquake M...	
1	4	NaN	NaN	Forest fire near La Ronge Sask. Canada	
2	5	NaN	NaN	All residents asked to 'shelter in place' are ...	
3	6	NaN	NaN	13,000 people receive #wildfires evacuation or...	
4	7	NaN	NaN	Just got sent this photo from Ruby #Alaska as ...	
5	8	NaN	NaN	#RockyFire Update => California Hwy. 20 closed...	
6	10	NaN	NaN	#flood #disaster Heavy rain causes flash flood...	
7	13	NaN	NaN	I'm on top of the hill and I can see a fire in...	
8	14	NaN	NaN	There's an emergency evacuation happening now ...	
9	15	NaN	NaN	I'm afraid that the tornado is coming to our a...	

	target	country	has_url	text_length	sentiment_score
0	1	NaN	False	69	0.63660
1	1	NaN	False	38	0.33000
2	1	NaN	False	133	0.35200
3	1	NaN	False	65	0.50000
4	1	NaN	False	88	0.50000
5	1	NaN	False	110	0.33000
6	1	NaN	False	95	0.18755
7	1	NaN	False	59	0.42345
8	1	NaN	False	79	0.30910
9	1	NaN	False	52	0.50000

okey, aunque me gustan las categoricas que puse, pareciera que están medio subrepresentadas en el dataset (pocos con country por ejemplo). Más adelante voy a revisar si tuvo sentido esto y si conviene cambiar features.

## 2.1 Model Training with GridSearchCV

Training logistic regression with hyperparameter search using F1 score.

```
[7]: from sklearn.model_selection import train_test_split, KFold
from scipy.sparse import hstack

# 1. Separar features y target
X = train_df[numeric_features + categorical_features + [embedding_feature]].
    copy()
y = train_df['target'].copy()

# 2. Split estratificado train/validation (80/20)
X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_SEED, stratify=y
)

print(f"Train set: {len(X_train)} samples")
print(f"Validation set: {len(X_val)} samples")
print(f"Target distribution in train: {y_train.value_counts().to_dict()}")
print(f"Target distribution in val: {y_val.value_counts().to_dict()}")
```

Train set: 6090 samples  
Validation set: 1523 samples  
Target distribution in train: {0: 3473, 1: 2617}  
Target distribution in val: {0: 869, 1: 654}

```
[8]: # Transformer para numéricas
scaler = StandardScaler()
X_train_numeric = scaler.fit_transform(X_train[numeric_features])
X_val_numeric = scaler.transform(X_val[numeric_features])
```

```

X_test_numeric = scaler.transform(test_df[numeric_features])

from sklearn.model_selection import KFold
from scipy.sparse import csr_matrix

# has_url -> OneHotEncoder
ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=True)

X_train_hasurl = ohe.fit_transform(
    X_train[['has_url']].fillna('missing')
)
X_val_hasurl = ohe.transform(
    X_val[['has_url']].fillna('missing')
)
X_test_hasurl = ohe.transform(
    test_df[['has_url']].fillna('missing')
)

print("\nOneHotEncoder (has_url):")
print(f" categories: {ohe.categories_[0].tolist()}")

```

```

OneHotEncoder (has_url):
categories: ['False', 'True']

```

```

[9]: # country -> Mean Target Encoding con K-Fold
country_col = 'country'
alpha = 10
global_mean = y_train.mean()

country_train = X_train[country_col].fillna('missing').astype(str)
country_val = X_val[country_col].fillna('missing').astype(str)
country_test = test_df[country_col].fillna('missing').astype(str)

kfold = KFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)

country_train_enc = pd.Series(index=country_train.index, dtype=float)

for train_idx, holdout_idx in kfold.split(country_train):
    # folds internos para evitar leakage
    ct_train = country_train.iloc[train_idx]
    y_fold = y_train.iloc[train_idx]

    stats = (
        pd.DataFrame({'country': ct_train, 'target': y_fold})
        .groupby('country')['target']
        .agg(['mean', 'count'])
    )

```

```

)
stats['smoothed'] = (
    stats['mean'] * stats['count'] + alpha * global_mean
) / (stats['count'] + alpha)

enc_values = country_train.iloc[holdout_idx].map(stats['smoothed'])
country_train_enc.iloc[holdout_idx] = enc_values

# categorías raras que no se mapearon en algún fold
country_train_enc = country_train_enc.fillna(global_mean)

# encoding para val/test: se calcula con TODO el train (ya no hay leakage)
full_stats = (
    pd.DataFrame({'country': country_train, 'target': y_train})
    .groupby('country')['target']
    .agg(['mean', 'count'])
)
full_stats['smoothed'] = (
    full_stats['mean'] * full_stats['count'] + alpha * global_mean
) / (full_stats['count'] + alpha)

country_val_enc = country_val.map(full_stats['smoothed']).fillna(global_mean)
country_test_enc = country_test.map(full_stats['smoothed']).fillna(global_mean)

print("\nMean encoding (country):")
print(f"  global_mean: {global_mean:.4f}")
print(f"  ejemplo valores train: {country_train_enc.head().to_dict()}")

# Pasar estos vectores a matrices sparse columna para combinarlos con hstack
X_train_country_enc = csr_matrix(country_train_enc.values.reshape(-1, 1))
X_val_country_enc = csr_matrix(country_val_enc.values.reshape(-1, 1))
X_test_country_enc = csr_matrix(country_test_enc.values.reshape(-1, 1))

# matriz categórica final = [has_url OHE, country_mean]
X_train_cat = hstack([X_train_hasurl, X_train_country_enc])
X_val_cat = hstack([X_val_hasurl, X_val_country_enc])
X_test_cat = hstack([X_test_hasurl, X_test_country_enc])

```

Mean encoding (country):  
 global\_mean: 0.4297  
 ejemplo valores train: {2721: 0.4156582304978332, 2259: 0.4156582304978332,  
 1815: 0.3698482572290349, 682: 0.42014738135427787, 7216: 0.4282758387941799}

[10]: # Transformer para texto (TF-IDF)  
 tfidf\_vectorizer = TfidfVectorizer(  
 max\_features=5000,

```

ngram_range=(1, 2),
min_df=2,
max_df=0.95,
strip_accents='unicode',
lowercase=True,
analyzer='word',
token_pattern=r'\w{1,}',
stop_words='english'
)
X_train_text = tfidf_vectorizer.fit_transform(X_train[embedding_feature] .
    ↪fillna(''))
X_val_text = tfidf_vectorizer.transform(X_val[embedding_feature].fillna(''))
X_test_text = tfidf_vectorizer.transform(test_df[embedding_feature].fillna(''))

print(f"\nTF-IDF vectorizer:")
print(f"  Vocabulary size: {len(tfidf_vectorizer.vocabulary_)}")
print(f"  Feature names (first 10): {tfidf_vectorizer.get_feature_names_out()[::
    ↪10].tolist()}")

```

TF-IDF vectorizer:  
Vocabulary size: 5000  
Feature names (first 10): ['0', '00', '00 http', '00 pm', '000', '01', '01  
04', '02', '03', '04']

[11]: # 4. Combinar todas las features

```

X_train_combined = hstack([X_train_numeric, X_train_cat, X_train_text])
X_val_combined = hstack([X_val_numeric, X_val_cat, X_val_text])
X_test_combined = hstack([X_test_numeric, X_test_cat, X_test_text])

print(f"\nCombined feature matrix:")
print(f"  Train shape: {X_train_combined.shape}")
print(f"  Validation shape: {X_val_combined.shape}")
print(f"  Test shape: {X_test_combined.shape}")

```

Combined feature matrix:  
Train shape: (6090, 5005)  
Validation shape: (1523, 5005)  
Test shape: (3263, 5005)

Okey ya tenemos nuestros sets separados. Ahora entrenamos el modelo con GridSearchCV para optimizar hiperparámetros. Busca todos los hiperparámetros en la grilla y usa validación cruzada para elegir los mejores según F1 score.

[15]:

```

from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, StratifiedKFold

```

```

from sklearn.linear_model import LogisticRegression
import numpy as np

param_grid = {
    'C': [0.1, 0.5, 1, 2, 5, 10, 20],
    'penalty': ['l2'],
    'solver': ['lbfgs', 'liblinear'],
    'class_weight': [None, 'balanced'],
    'max_iter': [500, 1000]
}

grid_search = GridSearchCV(
    LogisticRegression(),
    param_grid,
    cv=StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED),
    scoring='f1',
    n_jobs=-1,
    verbose=2
)

grid_search.fit(X_train_combined, y_train)

print("\nMejores parámetros:", grid_search.best_params_)
print(f"Mejor score F1 (CV): {grid_search.best_score_:.4f}")

best_model = grid_search.best_estimator_
y_val_pred = best_model.predict(X_val_combined)
val_f1 = f1_score(y_val, y_val_pred)

print(f"\n{'='*60}")
print(f"Score F1 en validación: {val_f1:.4f}")
print(f"{'='*60}")

```

Fitting 5 folds for each of 56 candidates, totalling 280 fits

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=liblinear;  
total time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=liblinear;  
total time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=liblinear;  
total time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=lbfgs; total  
time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=lbfgs; total  
time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=lbfgs; total  
time= 0.0s

[CV] END C=0.1, class\_weight=None, max\_iter=500, penalty=l2, solver=lbfgs; total  
time= 0.0s



```

[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=0.1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s

```

```
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
```

```

[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=0.5, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=lbfgs;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s

```

```

[CV] END C=1, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=1, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s

```

```

[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=12, solver=lbfsgs; total
time= 0.1s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=12, solver=lbfsgs; total
time= 0.1s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=2, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=5, class_weight=None, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s
[CV] END C=1, class_weight=balanced, max_iter=1000, penalty=12,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=12, solver=lbfsgs; total
time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=liblinear;
total time= 0.0s
[CV] END C=1, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=12, solver=lbfsgs;
total time= 0.1s

```

```

[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=l2,
solver=lbfgs; total time= 0.1s

[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s

```

```

[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=2, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=2, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=2, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s

```

```

[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s[CV] END C=10, class_weight=None, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s

[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s

```

```

[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=5, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.2s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=5, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=5, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s

```

```
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=liblinear;
total time= 0.0s
[CV] END C=10, class_weight=None, max_iter=500, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END C=20, class_weight=None, max_iter=1000, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2,
solver=liblinear; total time= 0.0s
```

```

[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=10, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=500, penalty=l2, solver=lbfgs;
total time= 0.2s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
[CV] END C=20, class_weight=balanced, max_iter=1000, penalty=l2, solver=lbfgs;
total time= 0.1s
Mejores parámetros: {'C': 2, 'class_weight': 'balanced', 'max_iter': 500,
'penalty': 'l2', 'solver': 'liblinear'}
Mejor score F1 (CV): 0.7494
=====
```

Score F1 en validación: 0.7615

=====

### 3 Conteste las preguntas:

3.1 ¿Cuál es el mejor score de validación obtenido? (¿Cómo conviene obtener el dataset para validar?)

```
[ ]: print(f"\n{'*60}'")
print(f"SCORE F1: {val_f1:.4f}")
print(f"{'*60}'")
```

```
=====
SCORE F1: 0.7615
=====
```

El dataset lo obtuvimos haciendo un split kfold del dataset original. El mejor F1 score de validación obtenido fue de aproximadamente 0.76.

### 3.2 ¿Cuáles son las features más importantes del modelo?

```
[17]: import matplotlib.pyplot as plt

coefficients = best_model.coef_[0]

# Crear plot
# Estructura: [numeric_features, has_url_False, has_url_True, country_encoded,
  tfidf_features]
feature_names = (
    numeric_features +
    ['has_url_False', 'has_url_True', 'country_encoded'] +
    list(tfidf_vectorizer.get_feature_names_out())
)

# Crear DataFrame con features y coeficientes
feature_importance = pd.DataFrame({
    'feature': feature_names,
    'coefficient': coefficients,
    'abs_coefficient': np.abs(coefficients)
}).sort_values('abs_coefficient', ascending=False)

print(f"Total features: {len(feature_names)}")
print(f"\nTop 30 features más importantes:")
print(feature_importance.head(30)[['feature', 'coefficient']].to_string())

# Graficar top 30
top_30 = feature_importance.head(30)

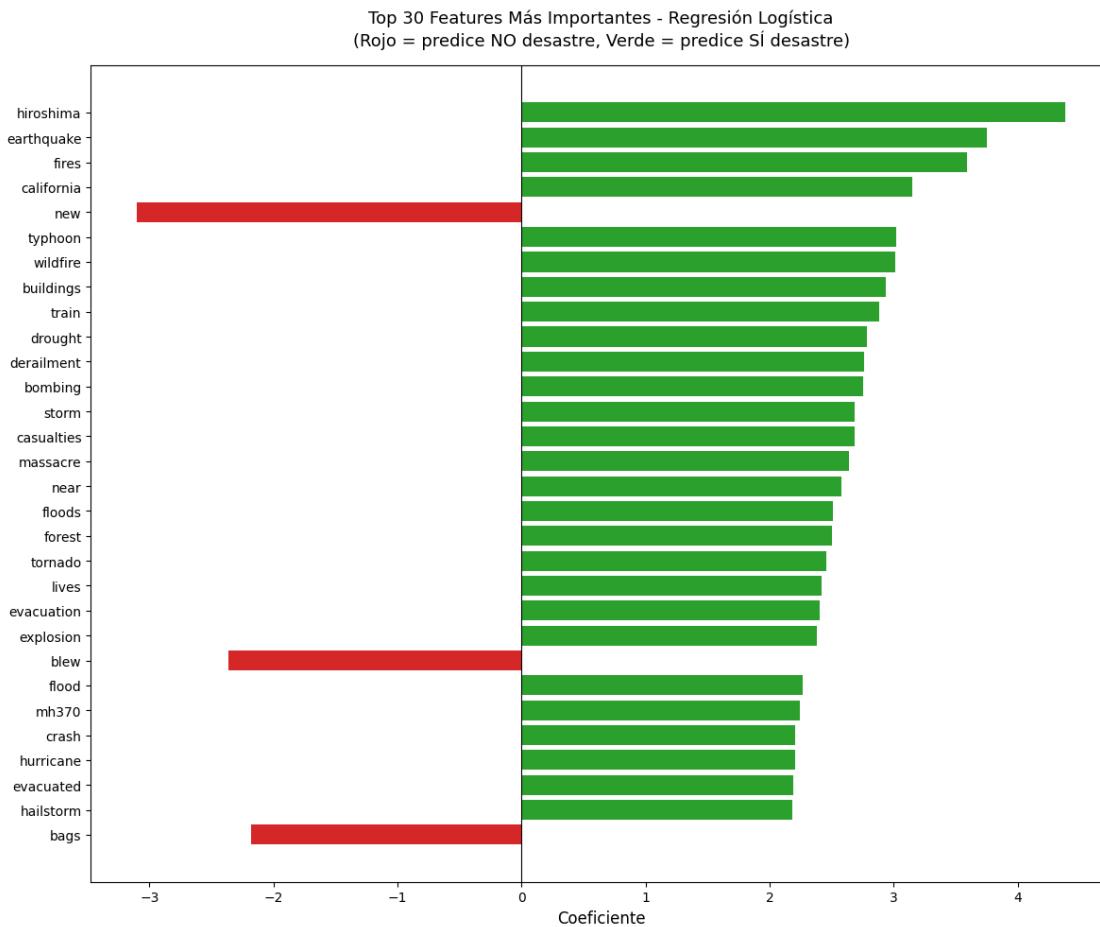
fig, ax = plt.subplots(figsize=(12, 10))
colors = ['#d62728' if c < 0 else '#2ca02c' for c in top_30['coefficient']]
ax.barh(range(len(top_30)), top_30['coefficient'], color=colors)
ax.set_yticks(range(len(top_30)))
ax.set_yticklabels(top_30['feature'], fontsize=10)
ax.axvline(x=0, color='black', linestyle='-', linewidth=0.8)
ax.set_xlabel('Coeficiente', fontsize=12)
ax.set_title('Top 30 Features Más Importantes - Regresión Logística\n(Rojo = predice NO desastre, Verde = predice SÍ desastre)', fontsize=13, pad=15)
ax.invert_yaxis()
plt.tight_layout()
plt.show()
```

Total features: 5005

Top 30 features más importantes:

	feature	coefficient
2151	hiroshima	4.381057

1393	earthquake	3.749714
1695	fires	3.593341
846	california	3.147560
3039	new	-3.099545
4606	typhoon	3.022060
4878	wildfire	3.014689
800	buildings	2.932460
4504	train	2.884047
1365	drought	2.783504
1244	derailment	2.763645
715	bombing	2.751401
4198	storm	2.685133
901	casualties	2.684507
2823	massacre	2.638483
3026	near	2.574520
1739	floods	2.506987
1775	forest	2.504547
4473	tornado	2.457713
2695	lives	2.415352
1491	evacuation	2.401101
1536	explosion	2.381438
665	blew	-2.365847
1730	flood	2.265254
2877	mh370	2.242496
1120	crash	2.205623
2269	hurricane	2.203791
1488	evacuated	2.187893
2022	hailstorm	2.182553
520	bags	-2.178557



Bueno, todas las features importantes salieron del TF IDF. Me hace sentir que perdí mi tiempo armando mis propias features... Pero bueno... Quizás igual sirvieron. Ninguna de las mías en el top 30 lamentablemente.

### 3.3 Al predecir con este modelo para la competencia, ¿Cúal es el score obtenido? (guardar el csv con predicciones para entregarlo después)

```
[19]: from pathlib import Path

submissions_dir = Path('../resultados')
submissions_dir.mkdir(parents=True, exist_ok=True)
baseline_submit_path = submissions_dir / 'logistic_regression_baseline.csv'

y_test_pred = best_model.predict(X_test_combined)

submission = pd.DataFrame({
    'id': test_df['id'],
    'desastre': y_test_pred
})
```

```

        'target': y_test_pred
    })

# Guardar CSV
submission.to_csv(baseline_submit_path, index=False)

print(f"\nArchivo guardado en: {baseline_submit_path}")
print(f"Total predicciones: {len(submission)}")
print(f"\nDistribución de predicciones:")
print(submission['target'].value_counts().sort_index())
print(f"\nPrimeras 10 filas:")
print(submission.head(10))

```

Archivo guardado en: ../resultados/logistic\_regression\_baseline.csv  
 Total predicciones: 3263

Distribución de predicciones:

```

target
0    1935
1    1328
Name: count, dtype: int64

```

Primeras 10 filas:

	id	target
0	0	1
1	2	1
2	3	1
3	9	0
4	11	1
5	12	1
6	21	0
7	22	0
8	27	0
9	29	0